

```
In [1]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
from ydata_profiling import ProfileReport
from sklearn.model_selection import train_test_split, GridSearchCV
from sklearn.tree import DecisionTreeClassifier, plot_tree
from sklearn.ensemble import RandomForestClassifier
from sklearn.metrics import accuracy_score, precision_score, recall_score, f1_score
from sklearn.preprocessing import LabelEncoder
from imblearn.over_sampling import SMOTE
from collections import Counter
import warnings
warnings.simplefilter(action = 'ignore')
```

```
In [2]: df = pd.read_csv(r"C:\Users\HP\Downloads\Indicino project.xlsx - Attrition_data.csv")
df.head()
```

Out[2]:

	Age	Attrition	BusinessTravel	DailyRate	Department	DistanceFromHome	Education
0	41	Yes	Travel_Rarely	1102	Sales	1	2
1	49	No	Travel_Frequently	279	Research & Development	8	1
2	37	Yes	Travel_Rarely	1373	Research & Development	2	2
3	33	No	Travel_Frequently	1392	Research & Development	3	4
4	27	No	Travel_Rarely	591	Research & Development	2	1

5 rows × 35 columns

```
In [3]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1470 entries, 0 to 1469
Data columns (total 35 columns):
 #   Column           Non-Null Count  Dtype  
--- 
 0   Age              1470 non-null    int64  
 1   Attrition        1470 non-null    object  
 2   BusinessTravel   1470 non-null    object  
 3   DailyRate        1470 non-null    int64  
 4   Department       1470 non-null    object  
 5   DistanceFromHome 1470 non-null    int64  
 6   Education        1470 non-null    int64  
 7   EducationField   1470 non-null    object  
 8   EmployeeCount    1470 non-null    int64  
 9   EmployeeNumber   1470 non-null    int64  
 10  EnvironmentSatisfaction 1470 non-null    int64  
 11  Gender            1470 non-null    object  
 12  HourlyRate       1470 non-null    int64  
 13  JobInvolvement   1470 non-null    object  
 14  JobLevel          1470 non-null    int64  
 15  JobRole           1470 non-null    object  
 16  JobSatisfaction  1470 non-null    int64  
 17  MaritalStatus     1470 non-null    object  
 18  MonthlyIncome     1470 non-null    int64  
 19  MonthlyRate       1470 non-null    int64  
 20  NumCompaniesWorked 1470 non-null    int64  
 21  Over18            1470 non-null    object  
 22  OverTime          1470 non-null    object  
 23  PercentSalaryHike 1470 non-null    int64  
 24  PerformanceRating 1470 non-null    int64  
 25  RelationshipSatisfaction 1470 non-null    int64  
 26  StandardHours     1470 non-null    int64  
 27  StockOptionLevel   1470 non-null    int64  
 28  TotalWorkingYears 1470 non-null    int64  
 29  TrainingTimesLastYear 1470 non-null    int64  
 30  WorkLifeBalance   1470 non-null    int64  
 31  YearsAtCompany    1470 non-null    int64  
 32  YearsInCurrentRole 1470 non-null    int64  
 33  YearsSinceLastPromotion 1470 non-null    int64  
 34  YearsWithCurrManager 1470 non-null    int64  
dtypes: int64(25), object(10)
memory usage: 402.1+ KB
```

```
In [4]: pip install ydata-profiling
```

```
Requirement already satisfied: ydata-profiling in c:\users\hp\anaconda3\lib\site-packages (4.12.2)
Requirement already satisfied: scipy<1.16,>=1.4.1 in c:\users\hp\anaconda3\lib\site-packages (from ydata-profiling) (1.13.1)
Requirement already satisfied: pandas!=1.4.0,<3,>1.1 in c:\users\hp\anaconda3\lib\site-packages (from ydata-profiling) (2.2.2)
Requirement already satisfied: matplotlib>=3.5 in c:\users\hp\anaconda3\lib\site-packages (from ydata-profiling) (3.9.2)
Requirement already satisfied: pydantic>=2 in c:\users\hp\anaconda3\lib\site-packages (from ydata-profiling) (2.8.2)
Requirement already satisfied: PyYAML<6.1,>=5.0.0 in c:\users\hp\anaconda3\lib\site-packages (from ydata-profiling) (6.0.1)
Requirement already satisfied: jinja2<3.2,>=2.11.1 in c:\users\hp\anaconda3\lib\site-packages (from ydata-profiling) (3.1.4)
Requirement already satisfied: visions<0.8.0,>=0.7.5 in c:\users\hp\anaconda3\lib\site-packages (from visions[type_image_path]<0.8.0,>=0.7.5->ydata-profiling) (0.7.6)
Requirement already satisfied: numpy<2.2,>=1.16.0 in c:\users\hp\anaconda3\lib\site-packages (from ydata-profiling) (1.26.4)
Requirement already satisfied: htmlmin==0.1.12 in c:\users\hp\anaconda3\lib\site-packages (from ydata-profiling) (0.1.12)
Requirement already satisfied: phik<0.13,>=0.11.1 in c:\users\hp\anaconda3\lib\site-packages (from ydata-profiling) (0.12.4)
Requirement already satisfied: requests<3,>=2.24.0 in c:\users\hp\anaconda3\lib\site-packages (from ydata-profiling) (2.32.3)
Requirement already satisfied: tqdm<5,>=4.48.2 in c:\users\hp\anaconda3\lib\site-packages (from ydata-profiling) (4.66.5)
Requirement already satisfied: seaborn<0.14,>=0.10.1 in c:\users\hp\anaconda3\lib\site-packages (from ydata-profiling) (0.13.2)
Requirement already satisfied: multimethod<2,>=1.4 in c:\users\hp\anaconda3\lib\site-packages (from ydata-profiling) (1.12)
Requirement already satisfied: statsmodels<1,>=0.13.2 in c:\users\hp\anaconda3\lib\site-packages (from ydata-profiling) (0.14.2)
Requirement already satisfied: typeguard<5,>=3 in c:\users\hp\anaconda3\lib\site-packages (from ydata-profiling) (4.4.1)
Requirement already satisfied: imagehash==4.3.1 in c:\users\hp\anaconda3\lib\site-packages (from ydata-profiling) (4.3.1)
Requirement already satisfied: wordcloud>=1.9.3 in c:\users\hp\anaconda3\lib\site-packages (from ydata-profiling) (1.9.4)
Requirement already satisfied: dacite>=1.8 in c:\users\hp\anaconda3\lib\site-packages (from ydata-profiling) (1.8.1)
Requirement already satisfied: PyWavelets in c:\users\hp\anaconda3\lib\site-packages (from imagehash==4.3.1->ydata-profiling) (1.7.0)
Requirement already satisfied: pillow in c:\users\hp\anaconda3\lib\site-packages (from imagehash==4.3.1->ydata-profiling) (10.4.0)
Requirement already satisfied: MarkupSafe>=2.0 in c:\users\hp\anaconda3\lib\site-packages (from jinja2<3.2,>=2.11.1->ydata-profiling) (2.1.3)
Requirement already satisfied: contourpy>=1.0.1 in c:\users\hp\anaconda3\lib\site-packages (from matplotlib>=3.5->ydata-profiling) (1.2.0)
Requirement already satisfied: cycler>=0.10 in c:\users\hp\anaconda3\lib\site-packages (from matplotlib>=3.5->ydata-profiling) (0.11.0)
Requirement already satisfied: fonttools>=4.22.0 in c:\users\hp\anaconda3\lib\site-packages (from matplotlib>=3.5->ydata-profiling) (4.51.0)
Requirement already satisfied: kiwisolver>=1.3.1 in c:\users\hp\anaconda3\lib\site-packages (from matplotlib>=3.5->ydata-profiling) (1.4.4)
Requirement already satisfied: packaging>=20.0 in c:\users\hp\anaconda3\lib\site-packages (from matplotlib>=3.5->ydata-profiling) (24.1)
```

```
Requirement already satisfied: pyparsing>=2.3.1 in c:\users\hp\anaconda3\lib\site-packages (from matplotlib>=3.5->ydata-profiling) (3.1.2)
Requirement already satisfied: python-dateutil>=2.7 in c:\users\hp\anaconda3\lib\site-packages (from matplotlib>=3.5->ydata-profiling) (2.9.0.post0)
Requirement already satisfied: pytz>=2020.1 in c:\users\hp\anaconda3\lib\site-packages (from pandas!=1.4.0,<3,>1.1->ydata-profiling) (2024.1)
Requirement already satisfied: tzdata>=2022.7 in c:\users\hp\anaconda3\lib\site-packages (from pandas!=1.4.0,<3,>1.1->ydata-profiling) (2023.3)
Requirement already satisfied: joblib>=0.14.1 in c:\users\hp\anaconda3\lib\site-packages (from phik<0.13,>=0.11.1->ydata-profiling) (1.4.2)
Requirement already satisfied: annotated-types>=0.4.0 in c:\users\hp\anaconda3\lib\site-packages (from pydantic>=2->ydata-profiling) (0.6.0)
Requirement already satisfied: pydantic-core==2.20.1 in c:\users\hp\anaconda3\lib\site-packages (from pydantic>=2->ydata-profiling) (2.20.1)
Requirement already satisfied: typing-extensions>=4.6.1 in c:\users\hp\anaconda3\lib\site-packages (from pydantic>=2->ydata-profiling) (4.11.0)
Requirement already satisfied: charset-normalizer<4,>=2 in c:\users\hp\anaconda3\lib\site-packages (from requests<3,>=2.24.0->ydata-profiling) (3.3.2)
Requirement already satisfied: idna<4,>=2.5 in c:\users\hp\anaconda3\lib\site-packages (from requests<3,>=2.24.0->ydata-profiling) (3.7)
Requirement already satisfied: urllib3<3,>=1.21.1 in c:\users\hp\anaconda3\lib\site-packages (from requests<3,>=2.24.0->ydata-profiling) (2.2.3)
Requirement already satisfied: certifi>=2017.4.17 in c:\users\hp\anaconda3\lib\site-packages (from requests<3,>=2.24.0->ydata-profiling) (2024.12.14)
Requirement already satisfied: patsy>=0.5.6 in c:\users\hp\anaconda3\lib\site-packages (from statsmodels<1,>=0.13.2->ydata-profiling) (0.5.6)
Requirement already satisfied: colorama in c:\users\hp\anaconda3\lib\site-packages (from tqdm<5,>=4.48.2->ydata-profiling) (0.4.6)
Requirement already satisfied: attrs>=19.3.0 in c:\users\hp\anaconda3\lib\site-packages (from visions<0.8.0,>=0.7.5->visions[type_image_path]<0.8.0,>=0.7.5->ydata-profiling) (23.1.0)
Requirement already satisfied: networkx>=2.4 in c:\users\hp\anaconda3\lib\site-packages (from visions<0.8.0,>=0.7.5->visions[type_image_path]<0.8.0,>=0.7.5->ydata-profiling) (3.3)
Requirement already satisfied: six in c:\users\hp\anaconda3\lib\site-packages (from patsy>=0.5.6->statsmodels<1,>=0.13.2->ydata-profiling) (1.16.0)
Note: you may need to restart the kernel to use updated packages.
```

```
In [5]: profile = ProfileReport(df, explorative=True)
profile.to_notebook_iframe()
```

```
Summarize dataset:  0% | 0/5 [00:00<?, ?it/s]
Generate report structure:  0% | 0/1 [00:00<?, ?it/s]
Render HTML:  0% | 0/1 [00:00<?, ?it/s]
```

Pandas Profiling Report



Overview

Brought to you by [YData](#)

[Overview](#)[Alerts 25](#)[Reproduction](#)

Dataset statistics

Number of variables 35**Number of observations** 1470**Missing cells** 0**Missing cells (%)** 0.0%**Duplicate rows** 0**Duplicate rows (%)** 0.0%**Total size in memory** 1.1 MiB**Average record size in memory** 770.8 B

Variable types

Numeric 15**Boolean** 3**Categorical** 17

Variables

[Select Columns](#)In [6]: `df.shape`

Out[6]: (1470, 35)

```
In [7]: # Handling outliers in numerical columns using z-scores method
# upon examining our dataset, we noticed that all numerical columns indicate no out
```

```
In [8]: category = df.select_dtypes(include = ['object', 'category']).columns
print(category)
```

```
Index(['Attrition', 'BusinessTravel', 'Department', 'EducationField', 'Gender',
       'JobInvolvement', 'JobRole', 'MaritalStatus', 'Over18', 'OverTime'],
      dtype='object')
```

```
In [9]: encoder = LabelEncoder()
df['Attrition'] = encoder.fit_transform(df['Attrition'])

categorical_cols = ['BusinessTravel', 'Department', 'EducationField', 'Gender', 'J
for cols in categorical_cols:
    df[cols] = encoder.fit_transform(df[cols])
```

```
In [10]: X = df.drop(columns = ['Attrition', 'StandardHours', 'Over18', 'EmployeeNumber', 'E
```

```
In [11]: y=df['Attrition']
```

```
In [12]: y.head()
```

```
Out[12]: 0    1
1    0
2    1
3    0
4    0
Name: Attrition, dtype: int32
```

```
In [13]: # (X = features, y = target variable)
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_st
```

```
In [14]: model = RandomForestClassifier(max_depth = 15, min_samples_split = 15, criterion =
model.fit(X_train, y_train)
```

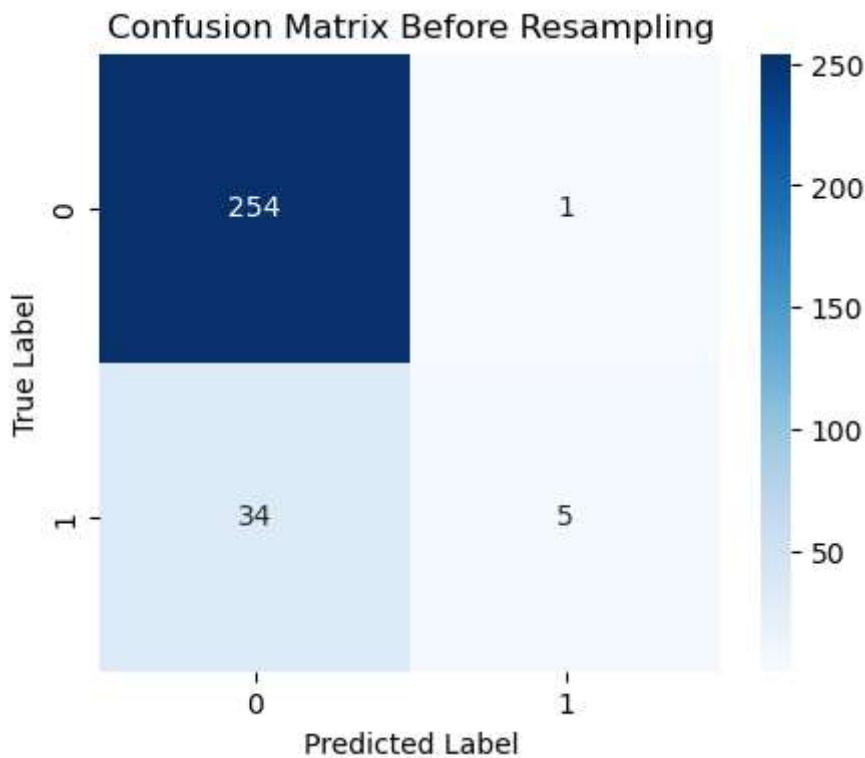
```
Out[14]: ▾ RandomForestClassifier
RandomForestClassifier(max_depth=15, min_samples_split=15, random_state=4
2)
```

```
In [15]: y_pred = model.predict(X_test)
```

```
In [37]: from sklearn.metrics import confusion_matrix
import seaborn as sns
import matplotlib.pyplot as plt
```

```
# Compute confusion matrix
cm = confusion_matrix(y_test, y_pred)

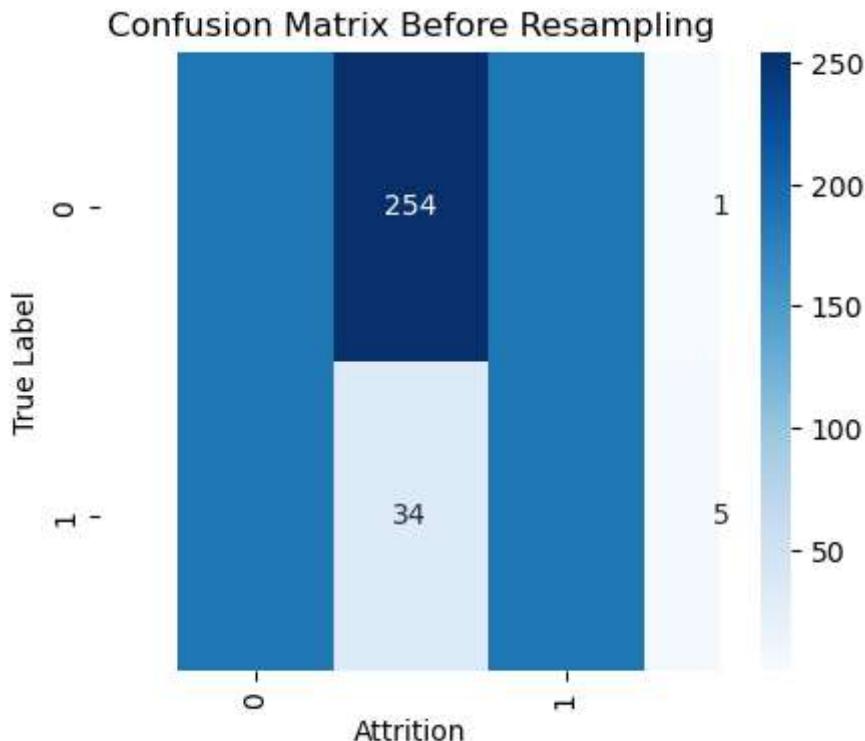
# Visualize it
plt.figure(figsize=(5, 4))
sns.heatmap(cm, annot=True, fmt="d", cmap="Blues")
plt.xlabel("Predicted Label")
plt.ylabel("True Label")
plt.title("Confusion Matrix Before Resampling")
plt.show()
```



In [17]:

```
%matplotlib inline
import matplotlib.pyplot as plt

y.value_counts().plot(kind='bar')
plt.show()
```



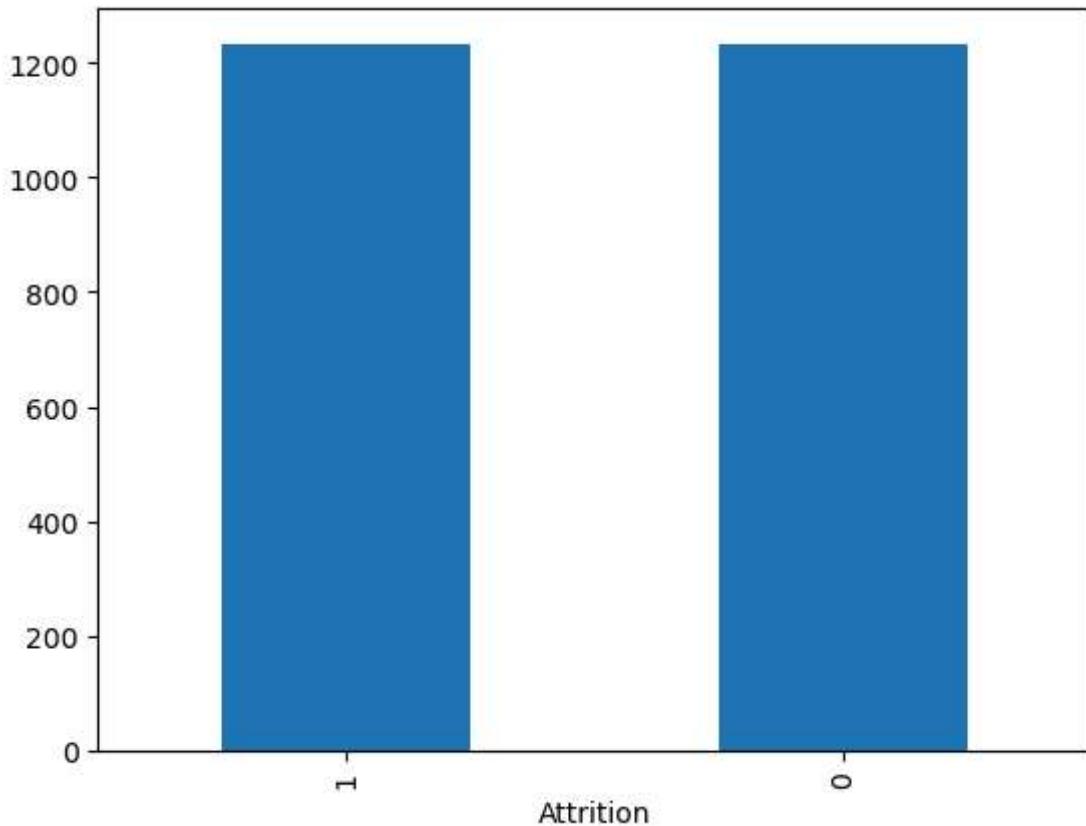
```
In [18]: # Apply SMOTE to balance the dataset
smote = SMOTE(random_state=42)

X_resampled, y_resampled = smote.fit_resample(X, y)

print("Before SMOTE:", y.value_counts()) # Class distribution before SMOTE
print("After SMOTE:", y_resampled.value_counts()) # Class distribution after S
```

```
Before SMOTE: Attrition
0    1233
1     237
Name: count, dtype: int64
After SMOTE: Attrition
1    1233
0    1233
Name: count, dtype: int64
```

```
In [19]: y_resampled.value_counts().plot(kind='bar')
plt.show()
```



```
In [20]: X_train, X_test, y_train, y_test = train_test_split(X_resampled, y_resampled, test
```

```
In [21]: model.fit(X_train, y_train)
```

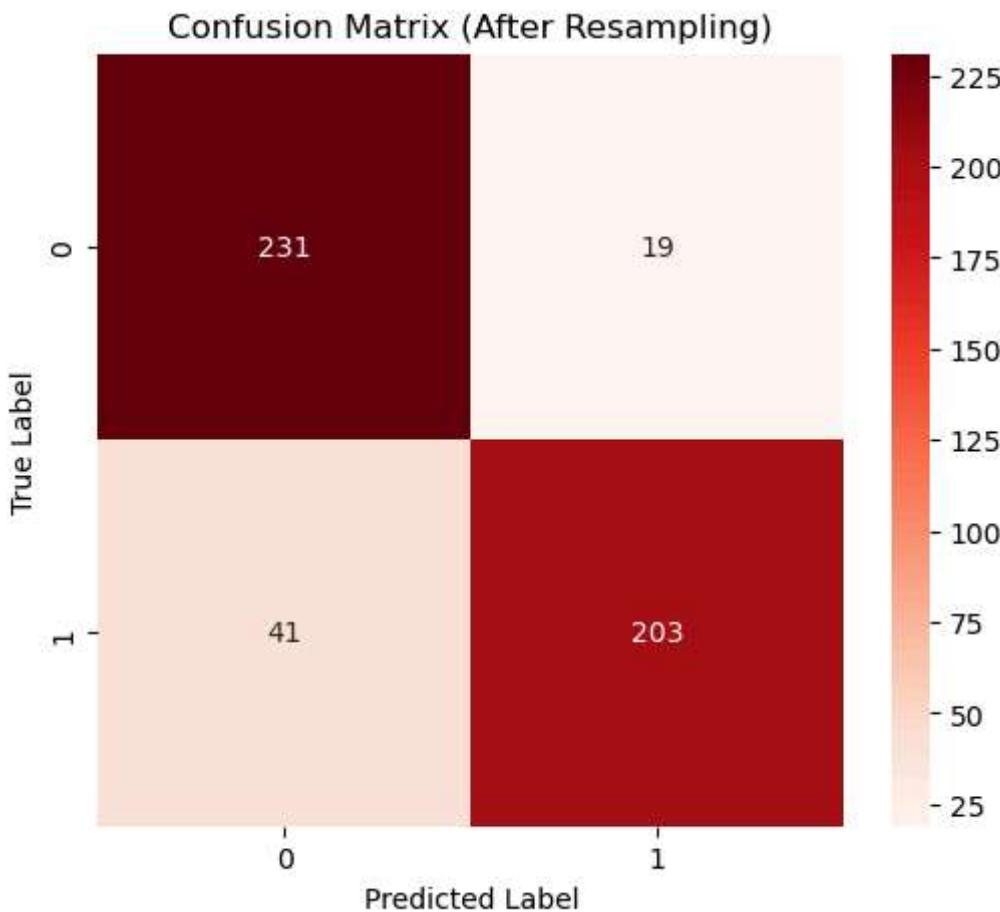
```
Out[21]: RandomForestClassifier
```

```
RandomForestClassifier(max_depth=15, min_samples_split=15, random_state=42)
```

```
In [22]: y_pred_resampled = model.predict(X_test)
```

```
In [23]: # Compute the confusion matrix
cm_resampled = confusion_matrix(y_test, y_pred_resampled)

# Plot the confusion matrix
plt.figure(figsize=(6, 5))
sns.heatmap(cm_resampled, annot=True, fmt="d", cmap="Reds", xticklabels=[0, 1], yti
plt.title('Confusion Matrix (After Resampling)')
plt.xlabel('Predicted Label')
plt.ylabel('True Label')
plt.show()
```



```
In [24]: print(f"y_pred shape: {y_pred.shape}")
```

```
y_pred shape: (294,)
```

```
In [25]: # Resampling changes only the training data (e.g., X_train_resampled, y_train_resampled)
# Some models, like Decision Trees, can still favor the majority class even after resampling
```

```
In [26]: print(y_test.shape, y_pred.shape)
```

```
(494,) (294,)
```

```
In [27]: # Classification Accuracy
```

```
from sklearn.model_selection import cross_val_score
from sklearn.metrics import accuracy_score, precision_score, recall_score, f1_score

# Cross-validation accuracy
scores = cross_val_score(model, X, y, cv=5, scoring="accuracy")
print(f'Cross-Validation Accuracy: {scores.mean():.2f}')

accuracy = accuracy_score(y_test, y_pred_resampled)
precision = precision_score(y_test, y_pred_resampled)
recall = recall_score(y_test, y_pred_resampled)
f1 = f1_score(y_test, y_pred_resampled)

# Compute ROC-AUC using predicted probabilities
y_pred_probs = model.predict_proba(X_test)[:, 1] # Get probabilities for class 1
```

```
roc_auc = roc_auc_score(y_test, y_pred_probs)

# Print results
print(f'Classification Accuracy: {accuracy:.2f}')
print(f'Precision: {precision:.2f}')
print(f'Recall: {recall:.2f}')
print(f'F1 Score: {f1:.2f}')
print(f'ROC_AUC: {roc_auc:.2f}')
```

Cross-Validation Accuracy: 0.86
Classification Accuracy: 0.88
Precision: 0.91
Recall: 0.83
F1 Score: 0.87
ROC_AUC: 0.95

```
In [28]: print("Unique values in y_resampled:", set(y_resampled))
```

Unique values in y_resampled: {0, 1}

```
In [29]: print("Type of y_test:", type(y_test))
print("Type of y_pred:", type(y_pred))
```

Type of y_test: <class 'pandas.core.series.Series'>
Type of y_pred: <class 'numpy.ndarray'>

```
In [30]: train_accuracy = model.score(X_train, y_train)

print(f"Train Accuracy: {train_accuracy:.2%}")
```

Train Accuracy: 97.77%

```
In [31]: print(f'Test Accuracy: {accuracy:.2%}')
```

Test Accuracy: 87.85%

```
In [32]: print(f'X_test shape: {X_test.shape}')
print(f'y_test shape: {y_test.shape}')
```

X_test shape:	Age	BusinessTravel	DailyRate	Department	DistanceFromHome	Ed
ucation \						
1078	44	2	136	1	28	3
1281	35	2	303	2	27	3
621	36	2	928	2	1	2
1508	29	1	336	1	15	1
1975	43	0	741	0	27	3
...
2099	47	2	1073	1	7	3
1836	29	0	477	1	23	2
303	31	2	218	2	7	3
1864	25	1	1123	1	3	3
141	45	2	1316	1	29	3
EducationField	EnvironmentSatisfaction	Gender	HourlyRate	\
1078	1	4	1	32	...	
1281	1	3	1	84	...	
621	1	2	1	56	...	
1508	3	3	0	81	...	
1975	0	1	0	77	...	
...
2099	4	2	0	62	...	
1836	1	4	1	73	...	
303	5	2	1	100	...	
1864	3	3	0	80	...	
141	3	3	1	83	...	
PerformanceRating	RelationshipSatisfaction	StockOptionLevel	\
1078	3	3	1			
1281	3	4	0			
621	3	4	1			
1508	3	1	0			
1975	3	3	1			
...
2099	3	4	0			
1836	4	2	0			
303	3	2	1			
1864	3	3	0			
141	3	2	0			
TotalWorkingYears	TrainingTimesLastYear	WorkLifeBalance	\
1078	24	1	4			
1281	10	2	3			
621	18	1	2			
1508	9	1	2			
1975	10	2	2			
...
2099	23	0	3			
1836	1	2	2			
303	10	3	2			
1864	4	4	3			
141	9	2	2			
YearsAtCompany	YearsInCurrentRole	YearsSinceLastPromotion	\
1078	20	6	14			
1281	10	7	7			

621	18	14	4
1508	7	6	6
1975	8	5	6
...
2099	14	9	8
1836	1	0	1
303	8	7	7
1864	3	1	0
141	6	5	0

YearsWithCurrManager	
1078	17
1281	7
621	11
1508	6
1975	7
...	...
2099	6
1836	0
303	7
1864	2
141	3

[494 rows x 30 columns]
y_test shape: 1078 0
1281 1
621 0
1508 1
1975 1
..
2099 1
1836 1
303 0
1864 1
141 0
Name: Attrition, Length: 494, dtype: int32

HYPERPARAMETER TUNING

```
In [34]: from sklearn.model_selection import GridSearchCV, RandomizedSearchCV
from sklearn.ensemble import RandomForestClassifier
from sklearn.metrics import classification_report

# Define the model
model = RandomForestClassifier(random_state=42)

# Define the hyperparameter grid
param_grid = {
    'n_estimators': [50, 100, 200],                      # Number of trees in the forest
    'max_depth': [None, 10, 20, 30],                      # Depth of each tree
    'min_samples_split': [2, 5, 10],                     # Minimum samples required to split a
    'min_samples_leaf': [1, 2, 4],                        # Minimum samples required at a Leaf n
    'max_features': ['sqrt', 'log2'],                    # Number of features to consider at ea
    'class_weight': ['balanced', None]                   # Handle class imbalance
```

```

}

# Initialize GridSearchCV (Exhaustive Search)
random_search = RandomizedSearchCV(model, param_grid, n_iter=50, cv=5, scoring='f1')
# Fit on training data
random_search.fit(X_resampled, y_resampled)

# Evaluate the best model
best_model = random_search.best_estimator_
y_pred = best_model.predict(X_test)

# Print classification report
print(classification_report(y_test, y_pred))

```

Fitting 5 folds for each of 50 candidates, totalling 250 fits

	precision	recall	f1-score	support
0	0.99	0.98	0.99	250
1	0.98	0.99	0.99	244
accuracy			0.99	494
macro avg	0.99	0.99	0.99	494
weighted avg	0.99	0.99	0.99	494

```
In [35]: print("Best Hyperparameters:", random_search.best_params_)
print("Best Score:", random_search.best_score_)
```

Best Hyperparameters: {'n_estimators': 200, 'min_samples_split': 2, 'min_samples_leaf': 1, 'max_features': 'log2', 'max_depth': 10, 'class_weight': None}
 Best Score: 0.8892839127223503

```
In [36]: best_model = random_search.best_estimator_
y_pred = best_model.predict(X_test)
```

```
In [37]: from sklearn.metrics import classification_report, confusion_matrix

print("Confusion Matrix:\n", confusion_matrix(y_test, y_pred))
print("Classification Report:\n", classification_report(y_test, y_pred))
```

Confusion Matrix:

[[245 5]	[2 242]]
----------	-----------

Classification Report:

	precision	recall	f1-score	support
0	0.99	0.98	0.99	250
1	0.98	0.99	0.99	244
accuracy			0.99	494
macro avg	0.99	0.99	0.99	494
weighted avg	0.99	0.99	0.99	494

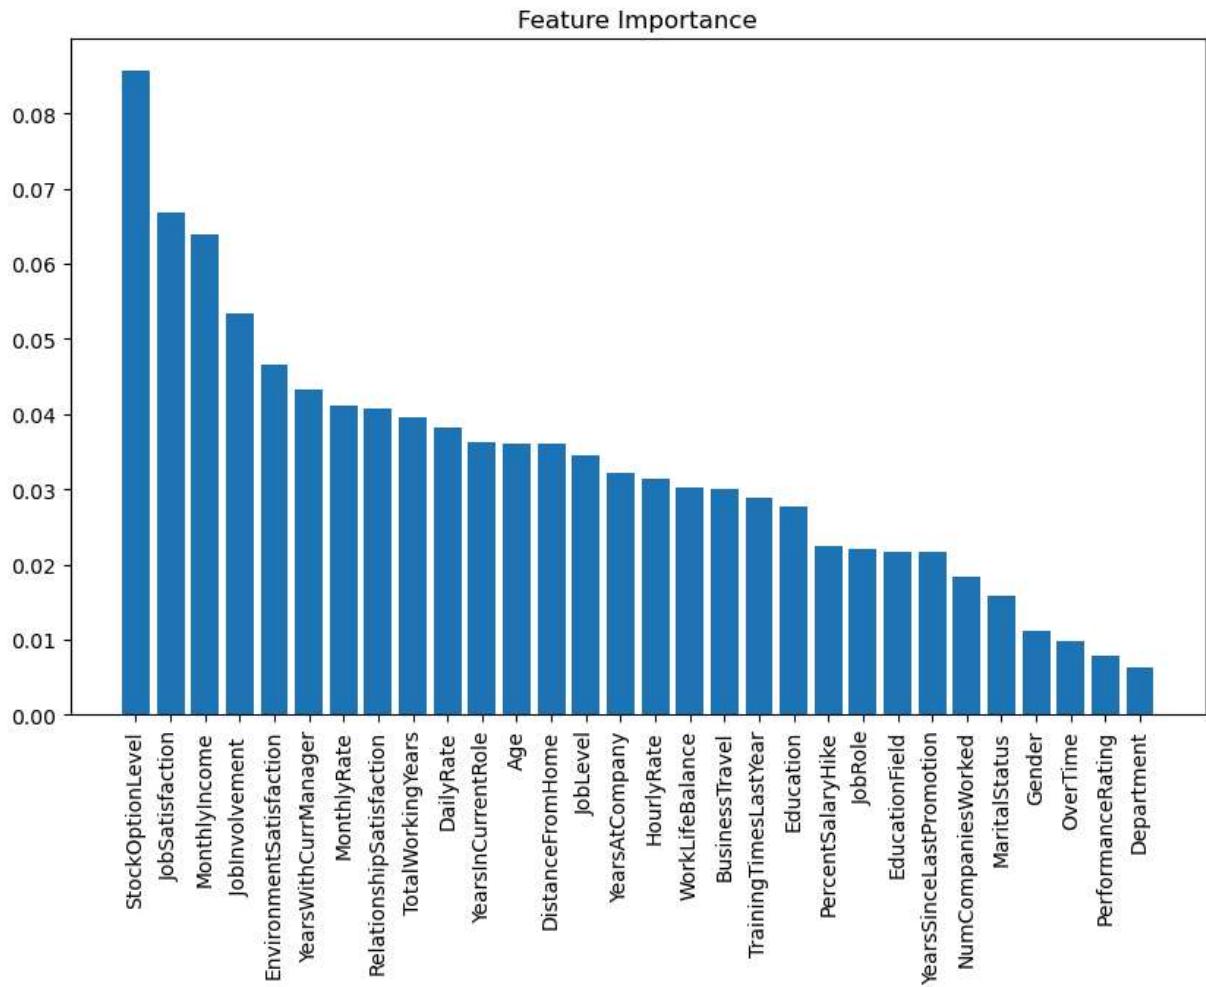
```
In [38]: import matplotlib.pyplot as plt
import numpy as np
```

```

feature_importances = best_model.feature_importances_
sorted_indices = np.argsort(feature_importances)[::-1] # Sort in descending order

plt.figure(figsize=(10, 6))
plt.bar(range(len(feature_importances)), feature_importances[sorted_indices])
plt.xticks(range(len(feature_importances)), X_train.columns[sorted_indices], rotation=90)
plt.title("Feature Importance")
plt.show()

```



```
In [39]: print("Train Accuracy:", best_model.score(X_train, y_train))
print("Test Accuracy:", best_model.score(X_test, y_test))
```

Train Accuracy: 0.9939148073022313
Test Accuracy: 0.9858299595141701

```
In [40]: pd.set_option('display.max_columns', None)
print(X.head())
```

	Age	BusinessTravel	DailyRate	Department	DistanceFromHome	Education	\
0	41	2	1102	2		1	2
1	49	1	279	1		8	1
2	37	2	1373	1		2	2
3	33	1	1392	1		3	4
4	27	2	591	1		2	1
	EducationField	EnvironmentSatisfaction	Gender	HourlyRate	\		
0	1		2	0	94		
1	1		3	1	61		
2	4		4	1	92		
3	1		4	0	56		
4	3		1	1	40		
	JobInvolvement	JobLevel	JobRole	JobSatisfaction	MaritalStatus	\	
0	2	2	7		4	2	
1	1	2	6		2	1	
2	1	1	2		3	2	
3	2	1	6		3	1	
4	2	1	2		2	1	
	MonthlyIncome	MonthlyRate	NumCompaniesWorked	Overtime	\		
0	5993	19479		8	1		
1	5130	24907		1	0		
2	2090	2396		6	1		
3	2909	23159		1	1		
4	3468	16632		9	0		
	PercentSalaryHike	PerformanceRating	RelationshipSatisfaction	\			
0	11		3		1		
1	23		4		4		
2	15		3		2		
3	11		3		3		
4	12		3		4		
	StockOptionLevel	TotalWorkingYears	TrainingTimesLastYear	\			
0	0		8		0		
1	1		10		3		
2	0		7		3		
3	0		8		3		
4	1		6		3		
	WorkLifeBalance	YearsAtCompany	YearsInCurrentRole	\			
0	1		6		4		
1	3		10		7		
2	3		0		0		
3	3		8		7		
4	3		2		2		
	YearsSinceLastPromotion	YearsWithCurrManager					
0	0			5			
1	1			7			
2	0			0			
3	3			0			
4	2			2			

```
In [95]: new_data = [[36, 2, 1973, 1, 2, 1, 4, 4, 1, 95, 1, 2, 3, 3, 3, 2190, 2596, 6, 4, 15
predictions = best_model.predict(new_data)
print("👤 Predictions:", predictions)

👤 Predictions: [0]
```

```
In [ ]:
```